

§6 — Oversight, Transparency & Contestability

Study Sheet 02 · the wheel, the label, the appeal · Guardrails 5 · 6 · 7

- Binding law (the floor) ■ Voluntary guidance ■ Frameworks & standards
- Controls (the toolbox) ■ Traps / red lines ■ Definitions


§6 of the Hawthorn Wealth AI Use Policy bundles three of the ten guardrails into one section. Each is a separate "beat" with its own rule.

One-line frame: §6 keeps a **human** — not the machine — in charge of the outcome at every stage: **before** it goes out (oversight), **while** it goes out (transparency), and **after**, if someone pushes back (contestability).

Beat 1 · The Wheel — Guardrail 5: Oversight

● Definition

Oversight = a human stays in control: positioned with the time, information and **authority to override** an AI output, and who **owns the outcome**.



● The trap

- Writing "**a human reviews AI outputs**" and thinking you're done. That's a **rubber stamp** — a glance and a signature.
- **Reviewing** = looked at it, signed it (passive). **Overseeing** = engaged, could override, owns it (active).

● The rule (scales with stakes)

- **High-stakes** (client advice, decisions) → reviewed + signed by a **named human** (the director).
- **Low-stakes** (internal drafts) → light check.
- **Why a regulator cares:** oversight puts a named human on the hook. The firm can't hide behind "the AI did it" — accountability needs **a human who could have said no**.

Beat 2 · The Label — Guardrail 6: Transparency

● Definition

Transparency = disclosing that AI was used, plainly, at the point it reaches the person — **where it matters to them**.

AI materially shaped it?

→

YES + affects client → TELL THEM |

trivial only → no label needed

● The trap (both directions)

- **Under:** burying "we may use AI" once in the fine print.
- **Over:** labelling every spell-check until it's meaningless noise.


● The rule — the hinge word is "materially"

- **Disclose** when AI materially shaped an output they receive or a decision about them (advice, plan, approval/denial).
- **Don't bother** for trivial backend work (spell-check, formatting, tidying a draft).
- "Materially" is what stops you having to label every keystroke. Transparency is about **the silence, not the AI** — people accept AI was used; they can't accept finding out you hid it.

Beat 3 · The Appeal — Guardrail 7: Contestability

● Definition

Contestability = a person affected by an AI-shaped outcome can **challenge** it, **get a human explanation, and have it reconsidered**. The escape hatch.



● The trap

- "We have a complaints inbox, so we're fine." An inbox only **receives** the objection — a black-box "no" nobody can explain or revisit is **not** contestable.

● The rule — three live parts (all must fire)

- **1. Told** — the person knows they can challenge it.
- **2. Human re-look + reason** — not the same AI; can explain why.
- **3. Real power to change** — overturn / adjust if the challenge is upheld.
- Contestability isn't only for when AI is wrong — it's the mechanism that lets wrong decisions **surface and be overturned** in the first place.

How it connects: §3 already **bans** "no human oversight" and "undisclosed AI on a client decision." §6 is where you say **what those protections actually are** — you cash the cheques §3 wrote. Oversight (before) and contestability (after) are the same idea at two ends: a human, not the machine, controls the outcome.

AI Governance study library · Sheet 02 · maps to **ai-use-policy.md §6** (Hawthorn Wealth Managers) ·
Source: Voluntary AI Safety Standard, Guardrails 5 · 6 · 7. Status: covered — needs drilling.